QANet.

Model

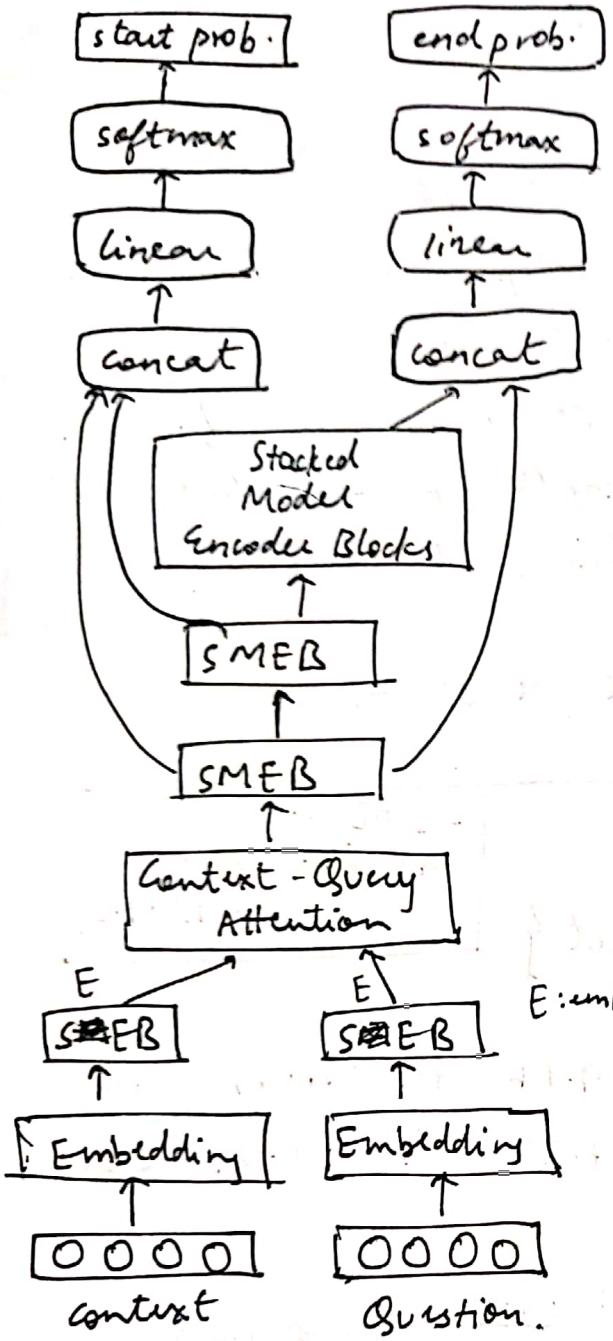

- 77.0 F1 score → 3hrs training (BiDAF: 15hrs)
- 82.7 F1 (dev) → 18hrs.

- SQuAD → 84.6 F1 (test) (81.8 SoTA 2017)
- 2018:

SQuAD	EM/F1
	82.2/88.6 (single)
	83.9/89.7 (ensemble)
	82.3 (human)

One Encoder Block.

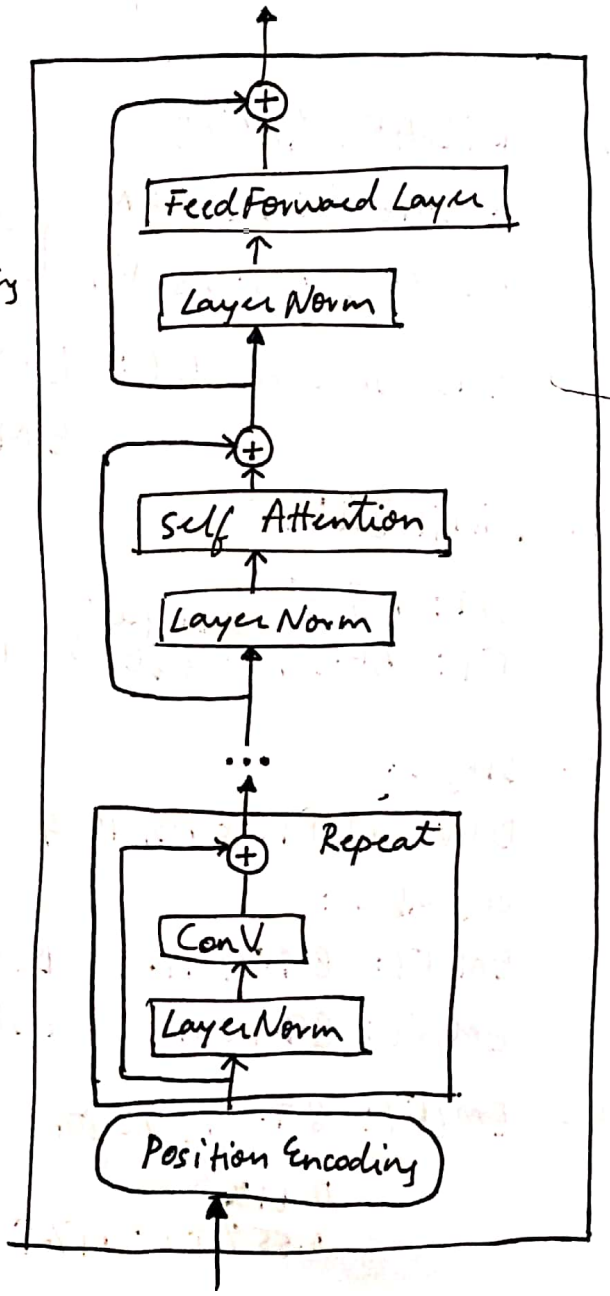

Model Layers:

1. Input Embedding L.
2. Embedding Encoder L.
3. Content-Query Attention L.
4. Model Encoder L.
5. Output L.

QANET

context paragraph : $C = \{c_1, \ldots, c_n\}$

query sentence : $Q = \{q_1, \ldots, q_m\}$

o/p span : $S = \{c_i, \ldots, c_{i+j}\}$

$x \in C, Q$

① **Input Embedding Layer:**

- word emb: fixed
  - $P_1 = 300$ D (GloVe)
  - <unk> → trainable & random init
- char emb:
  - $P_2 = 200$ D
  - max_word_len = 16
  - <pad>

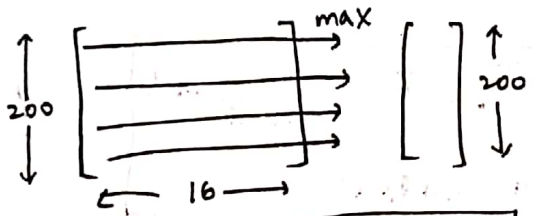

- $x = [x_w ; x_c] \in \mathbb{R}^{P_1+P_2} \in \mathbb{R}^{500}$

② **Embedding Encoder Layer**

- stack of the foll^n block:
  [conv layer x # + self attention layer + feed forward layer]
- depthwise separable conv. (mem. efficient)
  - $k = 7$
  - $d = 128$ filters
  - # conv layers within 1 block = 4.

  } conv.
- multiheaded attention
  - query, key ~~- inside a residual block.~~
  - no. of heads = 8

  } self atten.
- conv / attent / feed forw
  - inside a residual block
  - $f(\text{layernorm}(x) + x)$
  - no of encoder blocks = 1

  } General
- i/p = 500 D → 1D conv → 128 D *
  o/p = 128 D

  } i/p o/p

③ **Context-Query Attention Layer.**

- Similarity matrix $S \in \mathbb{R}^{n \times m}$: sim b/w C & Q. (n) (m)
- $S \xrightarrow{\text{softmax}} \bar{S}$ (row normalized) ⇒ dim = 1
- Attention (C→Q):

$$A = \bar{S} \cdot Q^T \in \mathbb{R}^{n \times d}$$
$$(n \times m)\ (m \times d)$$

- Similarity fn: (trilinear)

$$f(q, c) = W_0 [q, c, q \odot c]$$

* Query – Context Attention.

- (DCN)
- column normalized $S \xrightarrow{\text{soft}} \bar{\bar{S}}$ ⇒ dim = 0

$$B = \bar{S} \cdot \bar{\bar{S}}^T \cdot C^T \in \mathbb{R}^{n \times d}$$
$$(n \times m)\ (m \times n)\ (n \times d)$$

? $(m \times d)$

④ **Model Encoder Layer**

- i/p : $[c, a, c \odot a, c \odot b]$ a, b : row of A & B
- parameters same as before except:
  - #conv layers = 2 in a block
  - #blocks = 7

⇒ weights are shared b/w each of the 3 repitions of this Layer.

⑤ **Output Layer**

$$p^1 = \text{softmax}(W_1 [M_0; M_1]) \qquad p^2 = \text{softmax}(W_2 [M_0; M_2])$$

$$\text{score} = p^1 \times p^2$$

$$L(\theta) = -\frac{1}{N} \sum_{i}^{N} [\log(p^1_{y_i^1}) + \log(p^2_{y_i^2})]$$

Inference : $p_s^1 p_e^2 \rightarrow \max \quad s \le e$.

(DP can solve in linear time)

EXPERIMENTS
(SQuAD 1)

- **Data Preprocessing**
  - NLTK tokenizer
  - max context len = 400 (para longer → discard)
  - max ans len = 30
  - (300-D GLoVe)
  - <PAD> : short
  - <UNK> → trained
  - 200D char emb → trained

- **Training Details**
  - ① L2 weight decay
    - → all trainable variables
    - → $\lambda = 3 \times 10^{-7}$
  - ② Dropout
    - → word emb → 0.1
    - → char emb → 0.05
    - → blw layers → 0.1

- **Stochastic Depth Method (layer dropout)**
  - → within each embedding / model encoder layer.
  - → sublayer l : survival prob

$$p_l = 1 - \frac{l}{L}(1 - p_L)$$

$p_L = 0.9$ L = last layer.

- hidden size, conv filter } 128
- batch size = 32 steps = 150K
- conv : emb = 4 k = 7 # = 1
- conv : model = 2 k = 5 # = 7

- **ADAM**
  - $\beta_1 = 0.8$ $\varepsilon = 10^{-7}$
  - $\beta_2 = 0.999$
  - LR warm up scheme
    - 0.0 → 0.001 (1000 steps)
    - then constant.

- Exponential moving avg is applied on all trainable variables
  - decay rate = 0.9999.

NVIDIA P100 GPU.

	EM	F1
no data aug :	73.6	82.7

(SQuAD 1)